

1 Geometric Pattern Mining using MDL

We define geometric pattern mining on bounded, discrete and two-dimensional raster-based data. We represent this data as an $M \times N$ matrix A whose rows and columns are finite and in a fixed ordering (i.e. reordering rows and columns semantically alters the matrix). For elements $a_{i,j}$, where row i is on $[0; N)$ and column j is on $[0; M)$, holds that $a_{i,j} \in S$, the finite set of symbols over A .

According to MDL, the shortest (optimal) description of A reveals all structure of A in the most succinct way possible and we approximate this optimal description through the compression of the original data. This optimal description A' is only optimal if we can unambiguously reconstruct A from it and nothing more — the compression is both minimal and lossless. Intuitively compression means defining A using as few building blocks as possible. We illustrate this in Figure 1. Given the matrix A we decompose it in patterns, denoted X and Y . The set of all these patterns is the **model** of a A , denoted H_A . In order to reconstruct A from this model, we also need a mapping from the H_A back to A . This mapping represents in what MDL calls the **the data given the model** H_A . In this context we think of it as ‘instructions’ of how to reconstruct A . Let us call the set of all instructions required to rebuild A from H_A the **instantiation** of H_A . It is denoted by I_A in the example. The result is a notation that allows us to express matrix A as if decomposed into sets of local and global spatial information, which we will now describe in more detail.

1.1 Patterns and Instances

▷ We define a **pattern** as a $M_X \times N_X$ submatrix X of the original matrix A . Elements of this submatrix may be \cdot , the empty element, which gives us the ability to cut-out any irregular-shaped part of A . We additionally require the elements of X to be adjacent (horizontal, vertical or diagonal) to at least one non-empty element and that no rows and columns are empty.

From this definition, the dimensions $M_X \times N_X$ give the smallest rectangle around X (the *bounding box*). We also define the cardinality $|X|$ of X as the number of non-empty elements. We call a pattern X with $|X| = 1$ a **singleton pattern**, i.e. a pattern containing exactly one element of A .

Each pattern contains a special **pivot** element: $pivot(X)$ is the first non-empty element of X . A pivot can be thought of as a fixed point in X which we can use to position its elements in relation to A . This translation, or **offset**, is a tuple

$$A = \begin{bmatrix} 1 & \cdot & \cdot & 1 & 1 \\ \cdot & 1 & \cdot & \cdot & \cdot \\ 1 & \cdot & \cdot & 1 & \cdot \\ \cdot & 1 & \cdot & \cdot & 1 \\ 1 & 1 & 1 & 1 & \cdot \end{bmatrix}, I = \begin{bmatrix} X & \cdot & \cdot & Y & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ X & \cdot & \cdot & X & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ Y & \cdot & Y & \cdot & \cdot \end{bmatrix}, H = \{X = \begin{bmatrix} 1 & \cdot \\ \cdot & 1 \end{bmatrix}, Y = \begin{bmatrix} 1 & 1 \end{bmatrix}\}$$

Fig. 1: Example decomposition of A into instantiation I and patterns X, Y

$q = (i, j)$ that is on the same domain as an index in A . We realize this translation by placing all elements of X on an empty $M \times X$ size matrix in such that the pivot element is at (i, j) . We formalize this in the **instantiation operator** \otimes :

▷ We define the **instance** $X \otimes (i, j)$ as the $M \times N$ matrix containing all elements of X such that $\text{pivot}(X)$ is at index (i, j) and the distances between all elements are preserved. The resulting matrix contains no additional non-empty elements.

Obviously this does not yield a valid result for an arbitrary offset (i, j) . We want to limit ourselves to the space of pattern instances that are actually valid in relation to matrix A . Therefore two simple constraints are needed: (1) an instance must be **well-defined**: placing $\text{pivot}(X)$ is at index (i, j) results in an instance that contains all elements of X , and (2) elements of instances cannot overlap, meaning each element of A should be described at most once. This allows for a description that is both unambiguous and minimal.

▷ Two pattern instances $X \otimes q$ and $Y \otimes r$, with $q \neq r$ are **non-overlapping** if $|(X \otimes q) + (Y \otimes r)| = |X| + |Y|$.

From here on we will use the same letter in lower case to denote an arbitrary instance of a pattern, e.g. $x = X \otimes q$ when the exact value of q is unimportant. Since instances are simply patterns projected on an $M \times N$ matrix, we can reverse \otimes by removing all completely empty rows and columns:

▷ Let $X \otimes q$ be an instance of X , then by definition we say that $\odot(X \otimes q) = X$.

We briefly introduced the instantiation I as a set of ‘instructions’ of where instances of each pattern should be positioned in order to obtain A . As Figure 1 suggests, this mapping has the shape of an $M \times N$ matrix.

▷ Given the set of patterns H , the **instantiation (matrix)** I is an incomplete $M \times N$ matrix such that $I_{i,j} \in H \cup \{\cdot\}$ for all (i, j) . For all non-empty elements $I_{i,j}$ it holds that $I_{i,j} \otimes (i, j)$ is a non-overlapping instance of $I_{i,j}$ in A .

1.2 The Problem and its Solution Space

Patterns can be constructed by joining smaller patterns in a bottom-up fashion. We can define the exact way two patterns should be joined by enumerating the distance of their respective pivots. To limit the possibilities to patterns relevant to A , instances can be used as an intermediate step. As Figure 2 demonstrates, we can use a simple element-wise matrix addition to sum two instances and use \odot to obtain a joined pattern. Here we start by instantiating X and Y with offsets $(1, 0)$ and $(1, 1)$ respectively. We add the resulting x and y to obtain $\odot z$, the union of X and Y with relative offset $(1, 1) - (1, 0) = (0, 1)$.

The Sets \mathcal{H}_A and \mathcal{I}_A We define the **model class** \mathcal{H} , the set of all possible models for all possible inputs. Without any prior knowledge, this is the search space of our algorithm. We will first look at a more bounded subset \mathcal{H}_A of all possible models for A , and \mathcal{I}_A , the set of all possible instantiations to these models. We will also take H_A^0 to be the model with only singleton patterns. As

singletons are just individual elements of A , we can simply say that $H_A^0 = S$. The instantiation matrix corresponding to H_A^0 is denoted I_A^0 . Given that each element of this matrix must correspond to exactly one element of A in H_A^0 , we see that each $I_{i,j} = a_{i,j}$ and so I_A^0 is equal to A .

Using H_A^0 and I_A^0 as base cases we can now inductively define the set \mathcal{I}_A :

Base case $I_A^0 \in \mathcal{I}_A$

By induction If I is in \mathcal{I}_A then take any pair $I_{i,j}, I_{k,l} \in I$ such that $(i,j) \leq (k,l)$ in lexicographical order. Then the set I' is also in \mathcal{I}_A , providing I' equals I except:

$$\begin{aligned} I'_{i,j} &:= \oslash(I_{i,j} \otimes (i,j) + I_{k,l} \otimes (k,l)) \\ I'_{k,l} &:= \cdot \end{aligned}$$

This shows we can add any two instances together, which are by definition always non-overlapping and thus valid in A , in any order and obtain an element of \mathcal{I}_A . Eventually this results in just one big instance that is equal to A . Note that when we take two elements $I_{i,j}, I_{k,l} \in I$ we force $(i,j) \leq (k,l)$, not only to eliminate different routes to the same instance matrix, but also such that the pivot of the new pattern coincides with $I_{i,j}$. We can then leave $I_{k,l}$ empty.

The construction of \mathcal{I}_A also implicitly defines \mathcal{H}_A . While this may seem odd — defining models for instantiations instead of the other way around — note that there is no unambiguous way to find one instantiation for a given model. Instead we find the following definition by applying the inductive construction:

$$\mathcal{H}_A = \{ \{ \oslash(I) \mid I \in \mathcal{I}_A \} \mid I \in \mathcal{I}_A \}. \quad (1)$$

So for any instantiation $I \in \mathcal{I}_A$ there is a corresponding set in \mathcal{H}_A of all patterns that occur in I . This results in an interesting symbiosis between model and instantiation: increasing the complexity of one decreases that of the other. This construction gives a tightly connected lattice as shown in Figure 3.

1.3 Encoding Models and Instances

From all parametrized models in \mathcal{H}_A we want to select (approximate) the model that describes A best. We use two-part MDL to quantify how well a given model and instantiation matrix fit A . Two-part MDL tells us to minimize the sum of $L_1(H_A) + L_2(A|H_A)$, two functions that give the length of the model and the length of ‘the data given the model’, respectively. In this context, the model is the set of patterns H_A and the data given the model is I_A , the accidental information needed to reconstruct the data from H_A .

$$x = X \otimes (1, 0) = \begin{bmatrix} \cdot & \cdot \\ 1 & \cdot \\ \cdot & 1 \end{bmatrix}, \quad y = Y \otimes (1, 1) = \begin{bmatrix} \cdot & \cdot \\ \cdot & 1 \\ \cdot & \cdot \end{bmatrix}, \quad x + y = \begin{bmatrix} \cdot & \cdot \\ 1 & 1 \\ \cdot & 1 \end{bmatrix}, \quad Z = \oslash(x + y) = \begin{bmatrix} 1 & 1 \\ \cdot & 1 \end{bmatrix}$$

Fig. 2: Example of joining patterns X and Y to construct Z .




Fig. 3: Model space lattice for a 2×2 Boolean matrix. The Left hand column shows which pattern is added in each step, I is the current instantiation.

In order to compute their lengths, we need to decide on a way to encode H and I first. This encoding is of great influence on the length functions and therefore on the ability to accurately quantify the fit of a given H and I to A . Although the decision for this encoding is obviously prone to bias, practice shows that good results can be had once certain conditions are met: (1) all data is encoded (lossless) and (2) the encoding is as concise as possible (nothing but the data is encoded). Based on these conditions we give length functions for pattern sets and instantiations, but we do not actually need to encode them.

The fictional encoder sequentially sends each symbol in the datastream to the ‘decoder’ using a code word. Information theory tells us that the optimal length of a code word is given by $-\log(p)^1$, where p is the exact probability that the code word occurs in the output. We therefore need not compute the actual code words, just their probabilities. For this to work, both the encoder and hypothetical decoder must know either the exact probability distribution or agree upon an approximation beforehand. Such approximation is often called a **prior** and is used to fix ‘prior knowledge’ that does not have to be encoded explicitly.

The length function for incomplete matrices. To losslessly encode A' we have to encode both H and I individually. As both instances and patterns are both matrices it is tempting to utilize the same length function for both. Empirical evidence has shown that this is not a good idea though. For example, we do not consider certain values such as the size of the instance matrix because it is constant, however, the size of each individual pattern is not. We therefore have to construct a different length function for each type of matrix. These are listed in Table 1 When encoding I , we observe that it contains each pattern $X \in H$ multiple times, given by the **usage** of X . Using **prequential plug-in code** [?] to encode I enables us to omit encoding these usages separately, which would create unwanted bias. Prequential plug-in code gives us the following length

¹ We calculate lengths in bits and therefore all logarithms in this paper have base 2.

Table 1: Length computation for the different classes of matrices. The total length is the sum of the listed terms.

	Matrix	Bounds	# Elements	Positions	Symbols
$L_p(X)$	Pattern	$\log(MN)$	$L_N\binom{M_X N_X}{ X }$		$\log(S)$
$L_1(H)$	Model	N/A	$L_N(H)$	N/A	$L_p(X \in H)$
$L_2(I)$	Instantiation	<i>constant</i>	$\log(MN)$	<i>implicit</i>	$L_{pp}(I)$

function for I . We elaborate on the derivation of this equation in Appendix ??.

$$L_{pp}(I \mid P_{plugin}) = - \sum_{X_i \in h}^{|H|} \left[\log \frac{\Gamma(\text{usage}(X_i) + \epsilon)}{\Gamma(\epsilon)} \right] + \log \frac{\Gamma(|I| + \epsilon|H|)}{\Gamma(\epsilon|H|)} \quad (2)$$

Each length function has four terms. First we encode the total size of the matrix. Since we assume MN to be known/constant, we can use this constant to define the uniform distribution $\frac{1}{MN}$, such that $\log MN$ encodes an arbitrary index of A . Next we encode the number of elements that are non-empty. For patterns this value is encoded together with the third term, namely the positions of the non-empty elements. We use the previously encoded $M_X N_X$ in the binominal function to enumerate the ways we can place the $(|X|)$ elements onto a grid of $M_X N_X$. This gives us both *how many* non-empties there are as well as *where* they are. Finally the fourth term is the length of the actual symbols that encode the elements of matrix. In case we encode single elements of A , we simply assume that each unique value in A has an equal possibility of occurring. For the instance matrix, which encodes symbols to patterns, the prequential code is used as demonstrated before. Notice that L_N is the universal prior for integers[?] that can be used to encode integers on an arbitrary range.